



UNIVERSITÄT  
LEIPZIG

Medizinische Fakultät



Universitätsklinikum  
Leipzig

Medizin ist unsere Berufung.

# Machine Learning

## Introduction to Machine Learning

### Applications in Healthcare

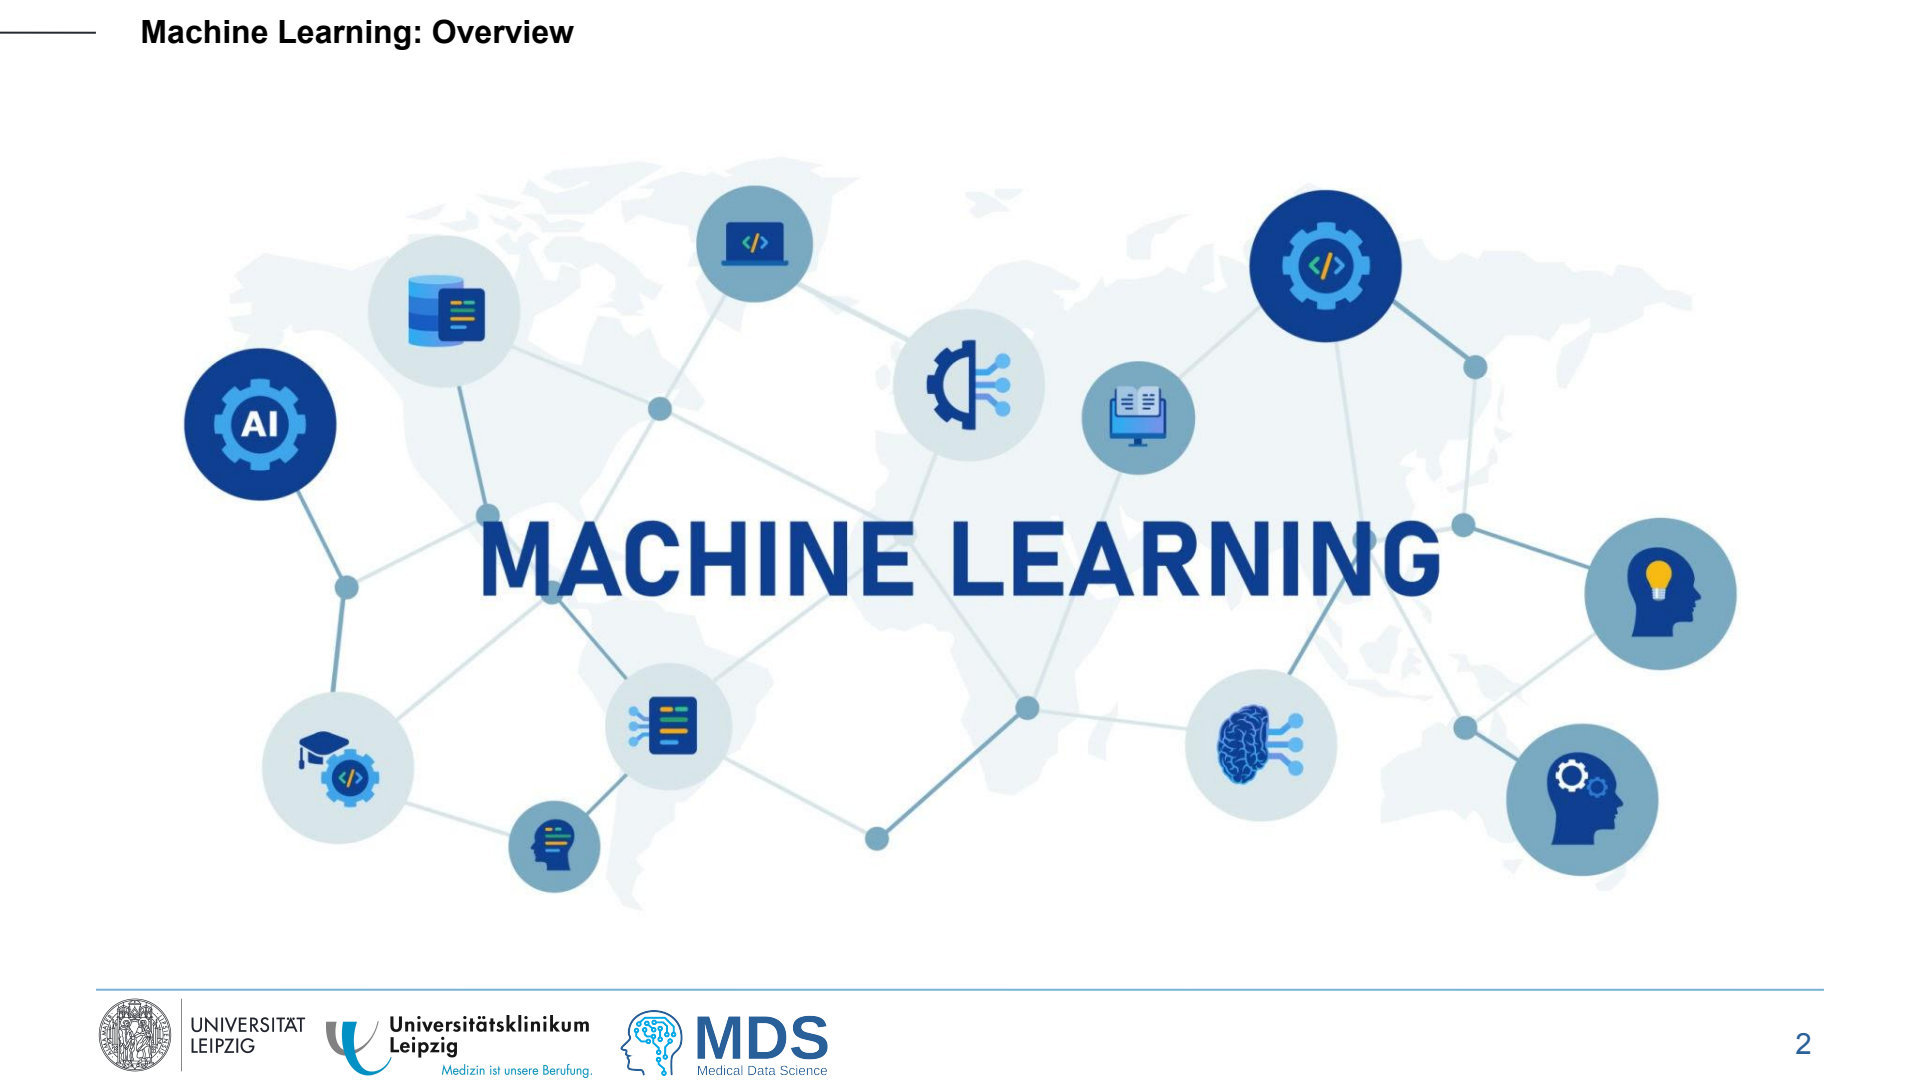
ScaDS.AI  
DRESDEN LEIPZIG

Data Science and AI for Medicine - Training School 2025

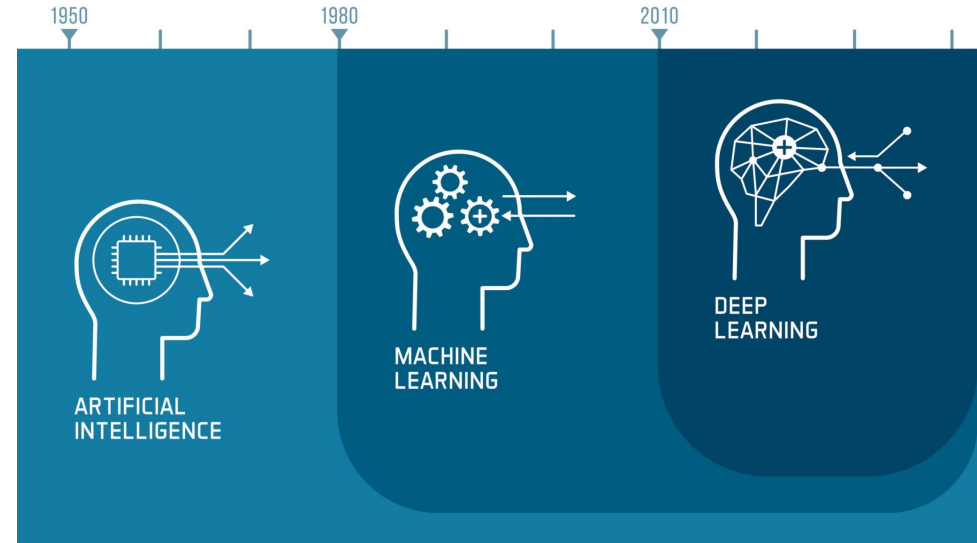



**MDS**  
Medical Data Science

Leipzig, 24 September 2025  
**Sina Sadeghi**




# Machine Learning vs Deep Learning & Artificial Intelligence



**AI:** techniques that enables computers to mimic human intelligence / behaviour

**ML:** ability of computers to learn without explicitly being programmed

**DL:** learning complex patterns from DATA using multi-layered artificial neural networks



### Applications in Medicine & Healthcare

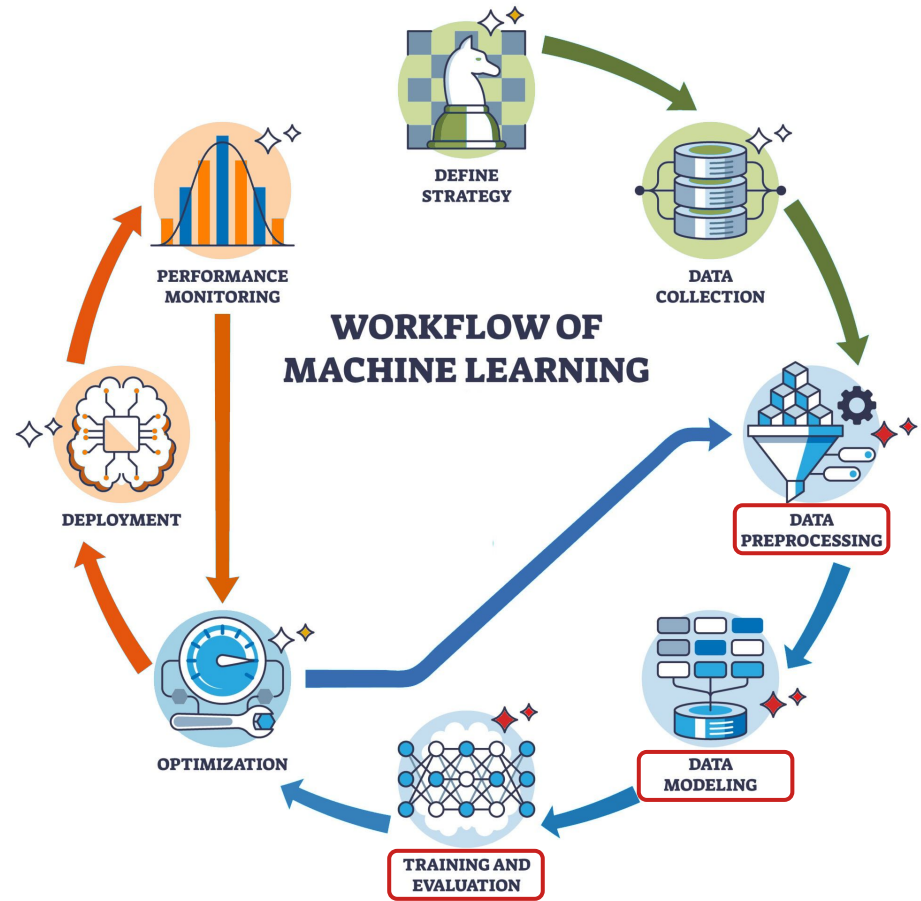
- ❖ **Diagnosis & Risk Prediction:** Detect diseases from images, labs, or genomics
- ❖ **Prognosis & Survival Analysis:** Predict disease progression, hospital stay
- ❖ **Treatment Recommendation:** Personalized medicine and drug dosing
- ❖ **Medical Imaging:** Detect tumors, segment organs, assist radiologists
- ❖ **Resource Optimization:** Predict patient flow, optimize hospital resources


**Arthur Samuel** (1959): Machine Learning is a field of study that gives computers the ability to learn without being explicitly programmed.

**Tom Mitchell** (1998): Well-posed Learning Problem – A computer program is said to learn from experience  $E$  with respect to some task  $T$  and some performance measure  $P$ , if its performance on  $T$ , as measured by  $P$ , improves with experience  $E$ .

## Define Machine Learning Problem


- Collect Data
- Preprocess Data
- Split Data (Train / Test)
- Select Model
- Train Model
- Tune Hyperparameters
- Evaluate Model






Supervised  
Learning

- ❖ Learns from **labeled data** (input → output)
- ❖ Goal: predict outcomes for new inputs
- ❖ Examples: classification, regression




Unsupervised  
Learning

- ❖ Learns patterns from **unlabeled data**
- ❖ Goal: discover structure or groupings
- ❖ Examples: clustering, dimensionality reduction




Reinforcement  
Learning

- ❖ Learns through interaction with an environment
- ❖ Goal: maximize cumulative reward via trial and error
- ❖ Examples: game playing, robotics




## Supervised Learning

- ❖ Learns from **labeled data** (input  $\rightarrow$  output)
- ❖ Goal: predict outcomes for new inputs
- ❖ Examples: classification, regression



## Unsupervised Learning

- ❖ Learns patterns from **unlabeled data**
- ❖ Goal: discover structure or groupings
- ❖ Examples: clustering, dimensionality reduction



## Reinforcement Learning


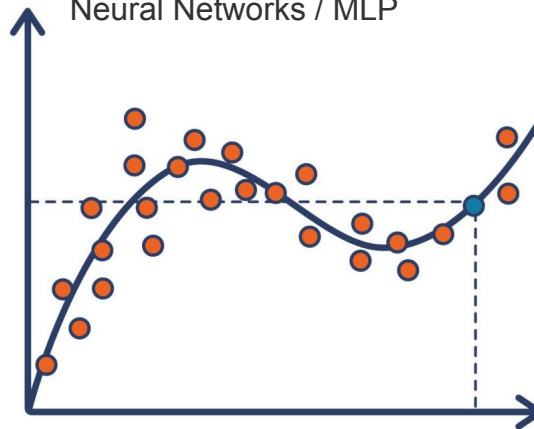
- ❖ Learns through interaction with an environment
- ❖ Goal: maximize cumulative reward via trial and error
- ❖ Examples: game playing, robotics



# Classification vs. Regression


## Regression

- ★ Predicts **continuous values**
- ★ Output: numerical values  
e.g., patient survival time (after diagnosis/surgery),  
hospital length of stay (for admitted patients)
- ★ Examples: Linear Regression,  
Neural Networks / MLP



## Classification

- ★ Predicts **discrete categories**
- ★ Output: class labels  
e.g., disease diagnosis: disease/no disease,  
classify tumors: benign vs. malignant
- ★ Examples: Logistic Regression, Decision Trees




### Classification

- ★ Predicts **discrete categories**
- ★ Output: class labels  
e.g., disease diagnosis: disease/no disease,  
classify tumors: benign vs. malignant
- ★ Examples: Logistic Regression, Decision Trees

### Regression

- ★ Predicts **continuous values**
- ★ Output: numerical values  
e.g., patient survival time (after diagnosis/surgery),  
hospital length of stay (for admitted patients)
- ★ Examples: Linear Regression,  
Neural Networks / MLP



## Dataset

## Pima Indians Diabetes Database

National Institute of Diabetes and Digestive and Kidney Diseases

feature 1

feature 2

↓

↓

$X_i$

$i = 1 \dots N$

ground truth

$y$

patient 1

→

patient 2

→

		Age	BMI	BloodPressure	Glucose	Insulin	SkinThickness	Pregnancies	DiabetesPedigree	Outcome
	0	50	33.6	72	148	0	35	6	0.627	1
	1	31	26.6	66	85	0	29	1	0.351	0
	2	32	23.3	64	183	0	0	8	0.672	1
	3	21	28.1	66	89	94	23	1	0.167	0
	4	33	43.1	40	137	168	35	0	2.288	1

patient<sub>j</sub>      $j = 1 \dots M$ 






Pima Indians Diabetes Database



Using the ADAP Learning Algorithm to Forecast the Onset of Diabetes Mellitus - PMC






# Classification

## Binary Classification






## Binary Classification



# Logistic Regression

input data



weights –  
model parameters

Classifier

$\hat{y}$

Binary Classification




1



0


Sigmoid function

$$g(z) = \frac{1}{1 + e^{-z}}$$



$$w_1 x_1 + w_2 x_2 + \dots = z$$

input data



weights –  
model parameters

Classifier

probability

$\hat{y}$

Binary Classification



1




0

Sigmoid function

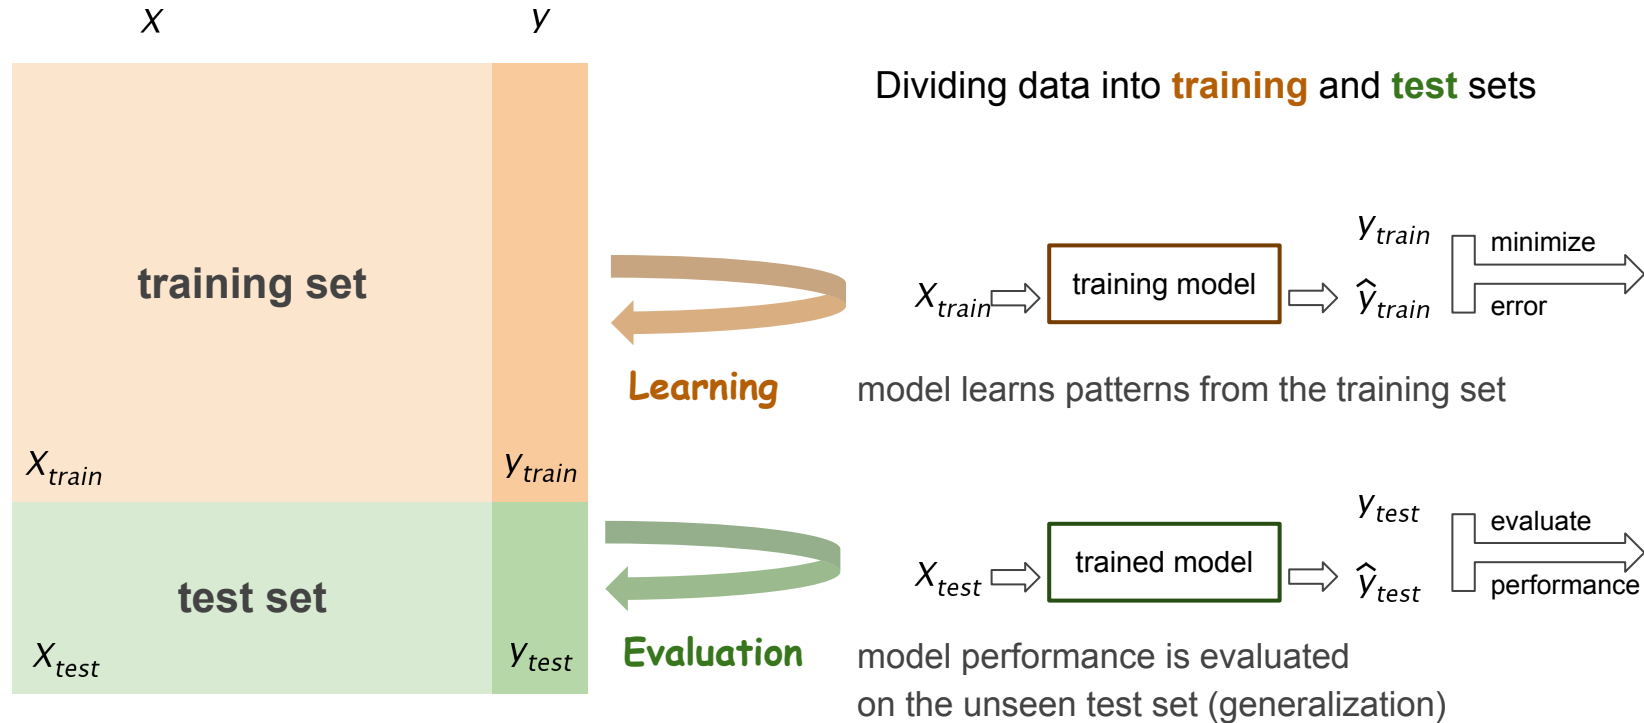
$$w_1x_1 + w_2x_2 + \dots = z$$

$$g(z) = \frac{1}{1 + e^{-z}}$$






# training vs. test



# Normalization

input data

Binary Classification



Rescaling numerical features to a common range (often [0, 1] or [-1, 1])

- ❑ Prevents features with large scales from dominating the model
- ❑ Improves convergence speed in optimization (e.g., gradient descent)

**Min-Max Scaling:**

$$x_{scaled} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

**Standardization:**



$$x_{scaled} = \frac{x - \mu}{\sigma}$$



$\mu$  mean

$\sigma$  standard deviation

# Optimization

**Goal:** minimizes the prediction error for binary classification





## Binary Classification



difference between what the model predicted and what the actual (true)

## Binary Cross Entropy

$M$  – number of instances

$\log$  – the natural logarithm

$y$  – binary true label (0 or 1)

$\hat{y}$  – predicted probability



$$\text{loss}(y, \hat{y}) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

$$\text{Loss}(y, \hat{y}) = \frac{1}{M} \sum_{i=1}^M \text{loss}(y_i, \hat{y}_i)$$

**Goal:** minimize Loss function

*Maths*

# Stochastic Gradient Descent



## Stochastic Gradient Descent (SGD)

an iterative optimization method that approximates gradient descent by using a small randomly selected subset of data

$$w_{new} = w_{old} - \alpha \frac{\partial \text{Loss}(y, \hat{y})}{\partial w}$$

$\alpha$  Learning rate

*Maths*

# Evaluation


## Performance metrics



- Measures how well a trained model performs on **unseen test data**
- Estimates how the model **generalizes** beyond the **training data**



# Classification – Evaluation Metrics



## Confusion Matrix

		Model Prediction	
		P	N
Ground Truth	P	TP	FN
	N	FP	TN

$$\text{Precision} = \frac{TP}{TP + FP}$$


$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

21	4
3	22

Acc = 0.86  
Sen = 0.84  
Spe = 0.88




0	2
1	47

Acc = 0.940  
Sen = 0.0  
Spe = 0.979

2	0
1	47

Acc = 0.98  
Sen = 1.0  
Spe = 0.979

# Classification – Evaluation Metrics



Confusion Matrix

		Model Prediction	
		P	N
Ground Truth	P	TP	FN
	N	FP	TN

$$\text{Precision} = \frac{TP}{TP + FP}$$


$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

21	4
3	22

Acc = 0.86  
Sen = 0.84  
Spe = 0.88



0	2
1	47


Acc = 0.940  
Sen = 0.0  
Spe = 0.979

2	0
1	47

Acc = 0.98  
Sen = 1.0  
Spe = 0.979



# Classification – Evaluation Metrics



Confusion Matrix

		Model Prediction	
		P	N
Ground Truth	P	TP	FN
	N	FP	TN

$$\text{Precision} = \frac{TP}{TP + FP}$$


$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

21	4
3	22

Acc = 0.86  
Sen = 0.84  
Spe = 0.88




0	2
1	47

Acc = 0.940  
Sen = 0.0  
Spe = 0.979

2	0
1	47

Acc = 0.98  
Sen = 1.0  
Spe = 0.979

# Classification – Evaluation Metrics



## Confusion Matrix

		Model Prediction	
		P	N
Ground Truth	P	TP	FN
	N	FP	TN

$$\text{Precision} = \frac{TP}{TP + FP}$$


$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

21	4
3	22

Acc = 0.86  
Sen = 0.84  
Spe = 0.88




0	2
1	47

Acc = 0.940  
Sen = 0.0  
Spe = 0.979

2	0
1	47

Acc = 0.98  
Sen = 1.0  
Spe = 0.979

# Hyperparameters



## Hyperparameters vs. Parameters in ML

### Hyperparameters


- Define the model's structure and training process
- Set before training (not learned from data)
- Examples: learning rate, number of layers, regularization strength

### Parameters

- Define how the model makes predictions
- Learned from data during training
- Examples: weights, biases in neural networks (and logistic regression)


# Under- vs. Overfitting

## Binary Classification




**Good fit** (robust model)

- ❑ Model captures patterns in data without excessive noise
- ❑ Performs well on both training and test dataset



**Underfitting** (simple model)

- ❑ Model is too simple, fails to capture patterns
- ❑ High error on training and test dataset



**Overfitting** (complex model)

- ❑ Model is too complex, memorizes training data
- ❑ Low training error but high test error

# Regularization

**Goal:** Prevents overfitting by penalizing large coefficients

**How:** Adds a penalty term to the loss function

❖ **L2 Regularization (Ridge)**

Shrinks coefficients smoothly (reduces their magnitude)

❖ **L1 Regularization (Lasso)**

Introduces sparsity (feature selection) – some coefficients become 0

**Effect:** Improves generalization and model robustness on unseen data

$$\text{Loss}(y, \hat{y}) = \mathcal{L}$$

$$\mathcal{L}_{L2} = \mathcal{L} + \lambda \sum_{j=1}^N w_j^2$$



$$\mathcal{L}_{L1} = \mathcal{L} + \lambda \sum_{j=1}^N |w_j|$$

$\lambda$  regularization strength (controls the penalty in the Loss function)

$\lambda = 0$  no regularization



*Maths*

## Binary Classification



# Learning

- ❑ Machine Learning Specialization
- ❑ AI for Medicine Specialization
- ❑ Machine Learning Mastery
- ❑ StatQuest with Josh Starmer - YouTube
- ❑ Big Data and Machine Learning in Health Care | Artificial Intelligence | JAMA
- ❑ Disease Prediction by Machine Learning Over Big Data From Healthcare Communities | IEEE Journals & Magazine





UNIVERSITÄT  
LEIPZIG

Medizinische Fakultät



Universitätsklinikum  
Leipzig

Medizin ist unsere Berufung.

**THANK YOU!**



**MDS**  
Medical Data Science

**Sina Sadeghi**  
Medical Data Science